

Databases

Lecture 10

Hash-Based Indexing

- hash function
 - maps search key values into a range of bucket numbers
- hashed file
 - search key (field(s) of the file)
 - pages grouped into *buckets*
 - determine record r's bucket
 - apply hash function to search key
 - quick location of records with given search key value
 - e.g., file hashed on *EmployeeName*
 - Find employee *Popescu*.
- ideal for equality selections

* hash functions

- relation R – key K
- records R stored in a file

$$\mathbf{F} = \{r_1, r_2, \dots, r_n\}$$

$$K_i = \Pi_K(r_i), i=1, \dots, n$$

- algorithm to determine the answer to the query: $K = K_0$
 - sequential search
 - examine an index

* hash functions

- define $h : \{K_1, K_2, \dots, K_n\} \rightarrow A$,
A = set of addresses that can store F
let $A = \{0, 1, \dots, m-1\}$, i.e., there are m locations that can store n records
- $h(K_i)$ = address where record r_i will be stored, $i=1, \dots, n$
- injectivity requirement
 - $h(K_i) \neq h(K_j), i \neq j$
 - dynamic collections - difficult
- efficiency – collisions are allowed
 - $h(K_i) = h(K_j), i \neq j$
 - r_i and r_j – synonyms
 - $h(K_i)$ - start address for search op.

- * hash functions

- functions that hash an integer value

- division method

- $h(K) = K \pmod{m} \Rightarrow$ range in $[0.. m-1]$

- e.g., $1618 \% 89 = 16$

- prime numbers - found to work best for m

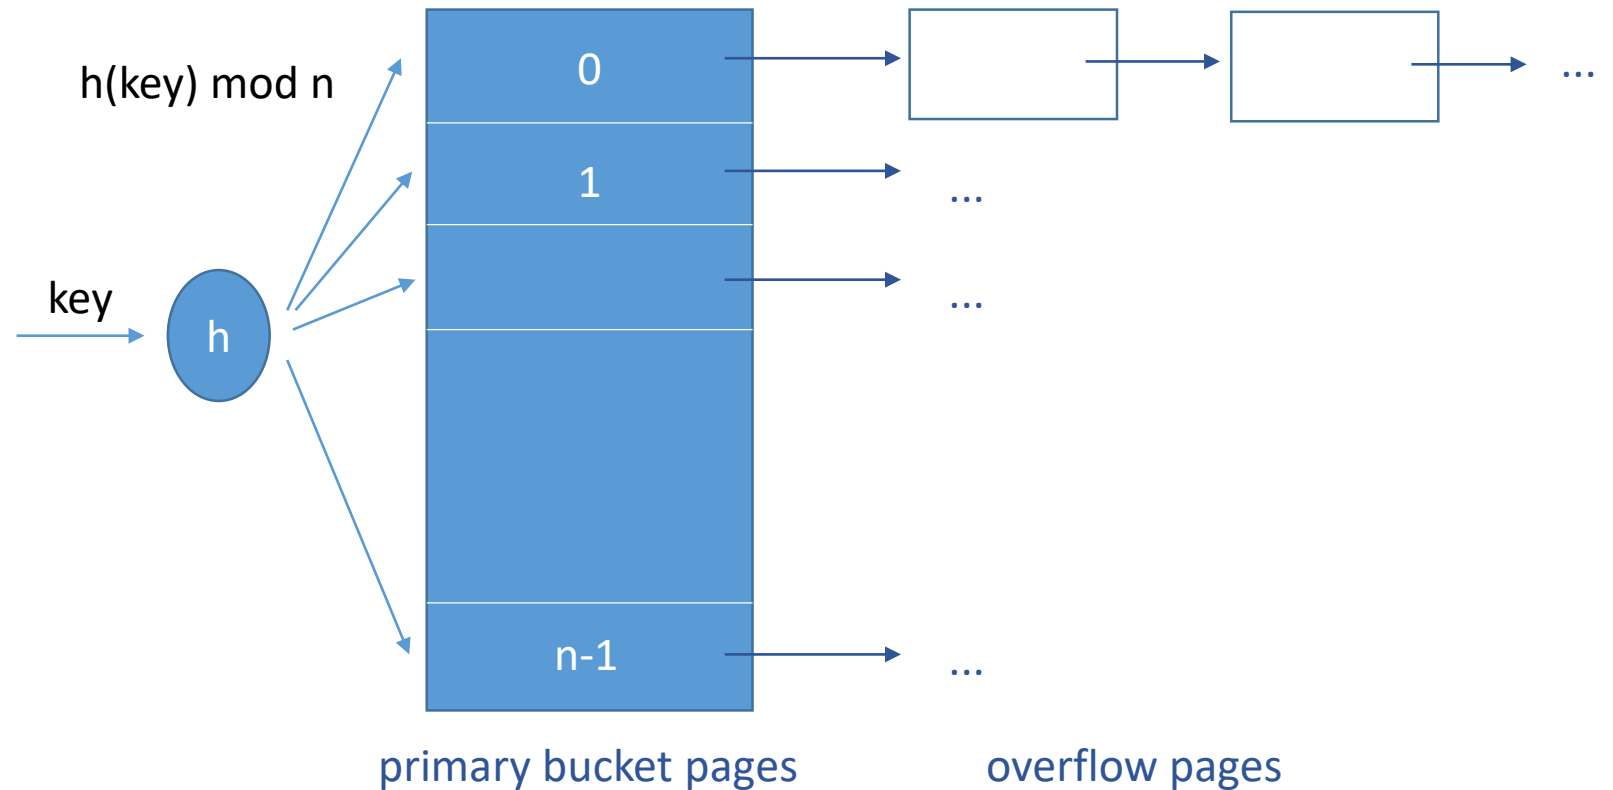
- multiplication method

- $h(K) = [m * \{Z * K\}]$

- good results for $Z = \frac{\sqrt{5}-1}{2} = 0.61803...$ or $1 - Z = 0.38196...$

- e.g., $[99 * \{0.61803 * 1618\}] = 96$

- * static hashing
 - buckets 0 to $n-1$
 - bucket
 - one primary page
 - possibly extra overflow pages
- data entries in buckets
 - $a_1/a_2/a_3$



* static hashing

- search for a data entry
 - apply hashing function to identify the bucket
 - search the bucket
 - possible optimization
 - entries sorted by search key
- add a data entry
 - apply hashing function to identify the bucket
 - add the entry to the bucket
 - if there is no space in the bucket:
 - allocate an overflow page
 - add the data entry to the page
 - add the overflow page to the bucket's overflow chain

* static hashing

- delete a data entry
 - apply hashing function to identify the bucket
 - search the bucket to locate the data entry
 - remove the entry from the bucket
 - if the data entry is the last one on its overflow page:
 - remove the overflow page from its overflow chain
 - add the page to a free pages list


- * static hashing
- good hashing function
 - few empty buckets
 - few records in the same bucket
 - i.e., key values are uniformly distributed over the set of buckets
 - good function in practice
 - $h(val) = a * val + b$
 - $h(val) \bmod n$ to identify bucket, for buckets numbered 0..n-1

- * static hashing
 - number of buckets known when the file is created
 - ideally
 - search = 1 I/O
 - I/D = 2 I/Os
 - file grows a lot => overflow chains; long chains can significantly affect performance
 - tackle overflow chains
 - initially, pages - 80% full
 - create a new file with more buckets
 - file shrinks => wasted space

- * static hashing
- main problem
 - fixed number of buckets
- solutions
 - periodic rehash
 - dynamic hashing


* extendible hashing

- dynamic hashing technique
- directory of pointers to buckets
- double the size of the number of buckets
 - double the directory
 - split overflowing bucket
- directory: array of 4 elements
- directory element: pointer to bucket
- entry r with key value K
- $h(K) = (... a_2 a_1 a_0)_2$
- $nr = a_1 a_0$, i.e., last 2 bits in $(... a_2 a_1 a_0)_2$, nr between 0 and 3
- $directory[nr]$: pointer to desired bucket



* extendible hashing

- global depth gd of hashed file
 - number of bits at the end of hashed value interpreted as an offset into the directory
 - kept in the header
 - depends on the size of the directory
 - e.g., 4 buckets $\Rightarrow gd = 2$
 - 8 buckets $\Rightarrow gd = 3$, etc
- initially, the global depth is equal to the local depth of every bucket




* extendible hashing

- insert entry

- find bucket

- a. bucket has free space =>
the new value can be added,
e.g., add data entry with hash
value 17 in bucket B

obs. data entry with hash value
17 is denoted as 17*



* extendible hashing

- insert entry


- b. bucket is full

- e.g., add entry 12*, bucket A full

- split bucket A

- allocate new bucket A'

- redistribute entries across A & A' (the split image of A), by taking into account the last 3 bits of $h(K)$




* extendible hashing

- insert entry

- b. bucket is full

- if $gd = \text{local depth of bucket}$ being split \Rightarrow double the directory, $gd++$
 - 3 bits are needed to discriminate between A & A', but the directory has only enough space to store numbers that can be represented on 2 bits, so it is doubled
 - increment local depth of bucket: $LD(A) = 3$
 - assign new local depth to bucket's split image: $LD(A') = 3$

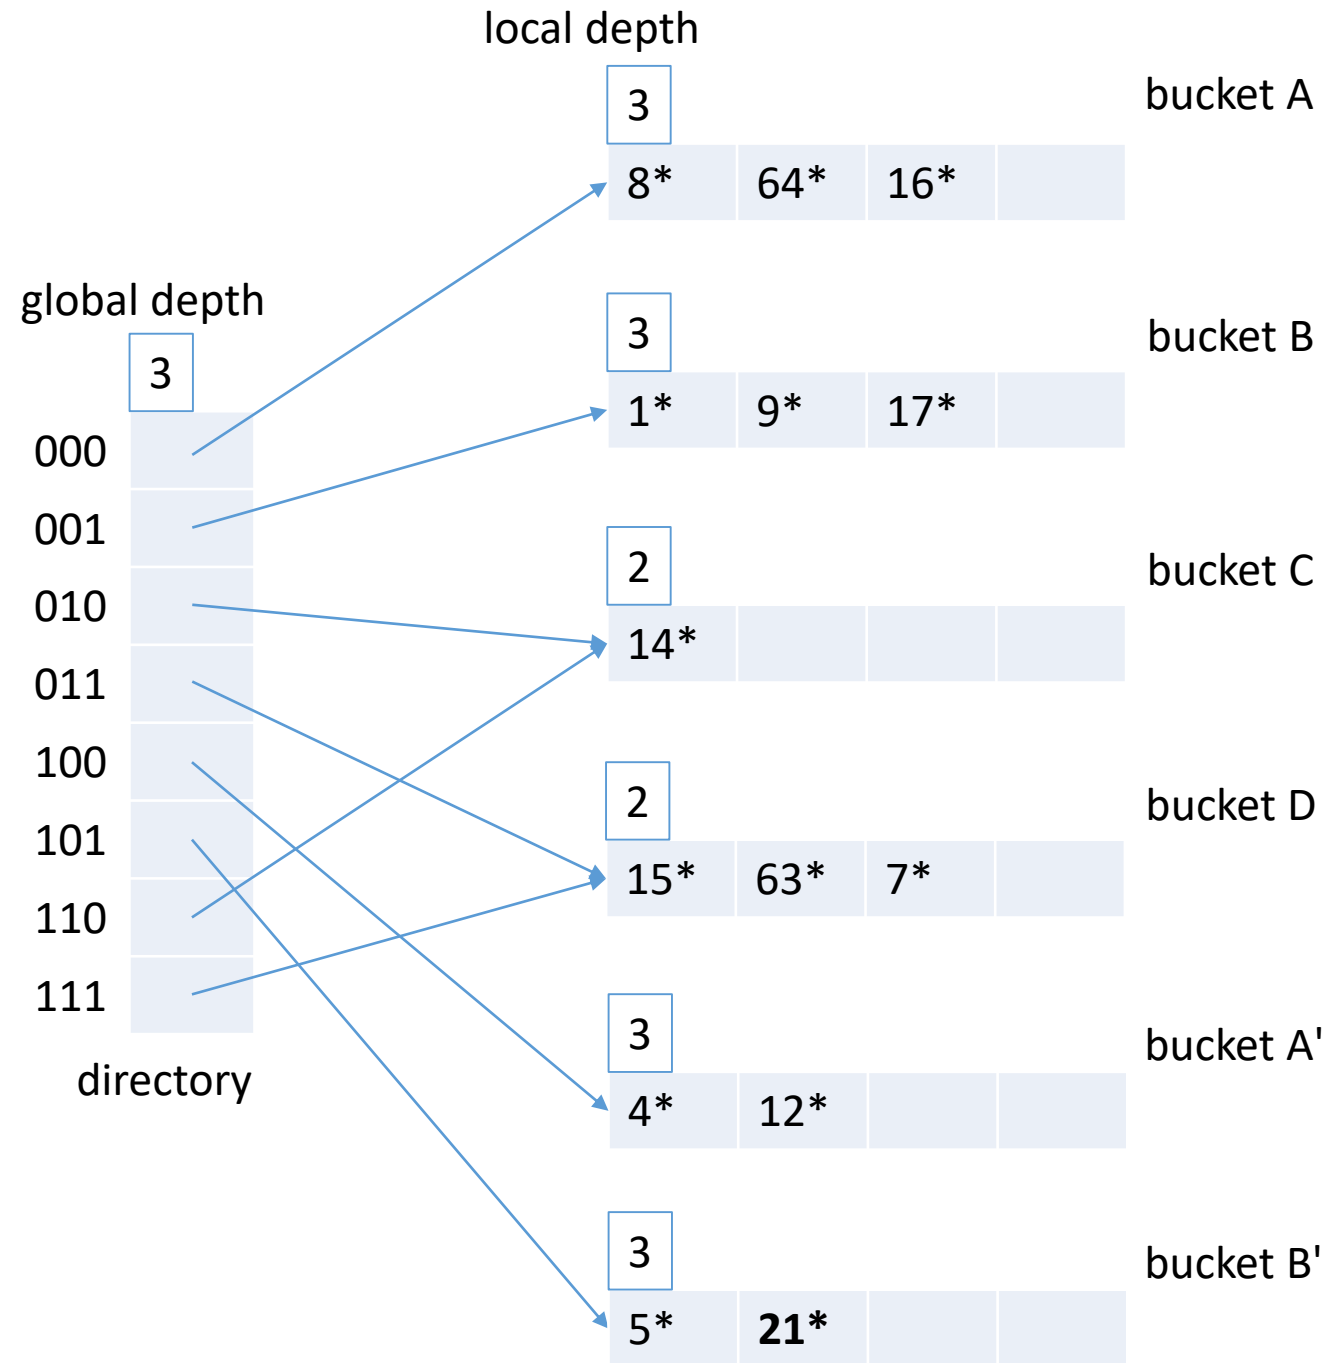


* extendible hashing

- insert entry

- b. bucket is full

- if $gd > \text{local depth of bucket}$ being split \Rightarrow directory doesn't have to be doubled
 - e.g., add 21^*
 - it belongs to bucket B, which is already full, but its local depth 2 is less than $gd = 3$
 - \Rightarrow split B, redistribute entries, increase local depth for B and its split image; directory isn't doubled, gd doesn't change



* extendible hashing

- search for entry with key value K_0
 - compute $h(K_0)$
 - take last gd bits to identify directory element
 - search corresponding bucket
- delete entry
 - locate & remove entry
 - if bucket is empty:
 - merge bucket with its split image, decrement local depth
 - if every directory element points to the same bucket as its split image:
 - halve the directory
 - decrement global depth

* extendible hashing

- obs 1. 2^{gd-l_d} elements point to a bucket Bk with local depth l_d
 - if $gd=l_d$ and bucket Bk is split => double directory
- obs 2. manage collisions - overflow pages
- double number of buckets in static file to avoid overflow pages
 - read entire file
 - write twice as many pages
- double extendible hashed file
 - allocate new bucket page nBk
 - write nBk and its split image
 - double directory array (which should be much smaller than file, since it has 1 page-id / element)
 - if using *least significant bits* (last gd bits) => efficient operation:
 - copy directory over
 - adjust split buckets' elements

- * extendible hashing
 - equality selection
 - if directory fits in memory:
 - => 1 I/O (as for Static Hashing with no overflow chains)
 - otherwise
 - 2 I/Os
 - e.g., 100 MB file, entry = 50 bytes => 2.000.000 entries
 - page size = 8 KB => approx. 160 entries / bucket
- => need $2.000.000 / 160 = 12.500$ directory elements

References

- [Ra00] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (2nd Edition), McGraw-Hill, 2000
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007,
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Ta13] ȚÂMBULEA, L., Curs Baze de date, Facultatea de Matematică și Informatică, UBB, 2013-2014
- [Si10] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, McGraw-Hill, 2010
- [Ga08] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book, Prentice Hall Press, 2008