

# Databases

## Seminar 5

### Indexes in SQL Server (I)

- index
  - structure stored on the disk, associated with a table or view
  - optimizes retrieval operations on the table / view
- great indexing
  - => fast applications
- poor indexing
  - => can slow down the DBMS

- syntax

```
CREATE [ UNIQUE ] [ CLUSTERED | NONCLUSTERED ]  
    INDEX index_name  
  
ON <object> ( column [ ASC | DESC ] [ ,...n ] )  
  
[ INCLUDE ( column_name [ ,...n ] ) ]  
  
[ WHERE <filter_predicate> ]  
  
[ WITH ( <index_option> [ ,...n ] ) ]
```

- index characteristics
  - clustered versus non-clustered
  - unique versus non-unique
  - search key - single-column versus multicolumn
  - key columns versus non-key columns
  - columns in the index - ascending / descending order
  - non-clustered indexes - full-table versus filtered

- clustered versus non-clustered index
  - clustered index
    - the data rows in the table are kept sorted, based on the values of the search key

```
CREATE CLUSTERED INDEX Index_Name  
    ON Schema_Name.Table_Name(Column)
```

- non-clustered index
  - contains key values and pointers to the data in the table (heap / clustered index)

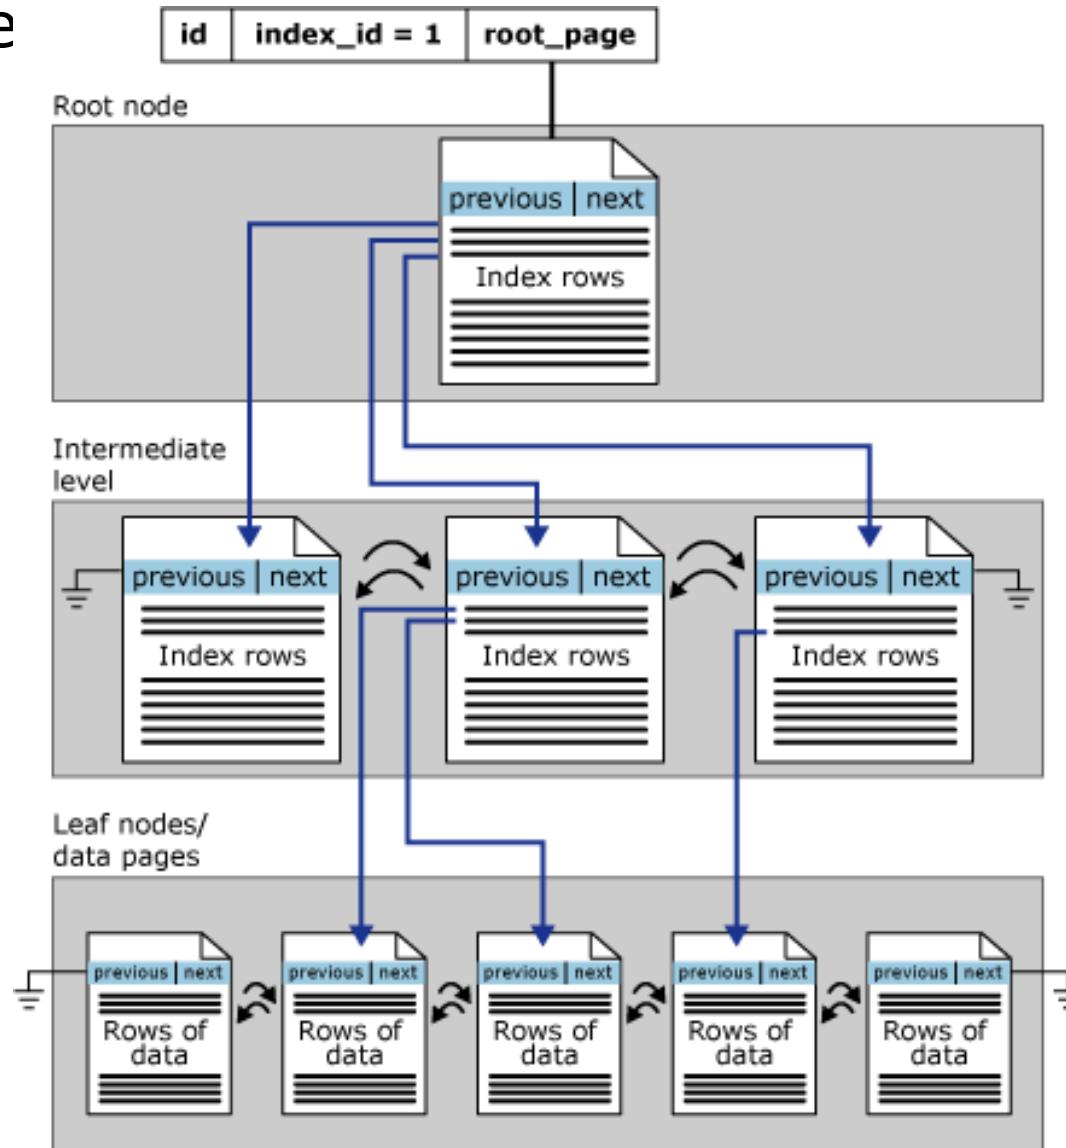
```
CREATE INDEX Index_Name  
    ON Schema_Name.Table_Name(Column)
```

- clustered versus non-clustered index
  - data pages in a clustered index always include *all the columns* in the table
  - SQL Server:
    - at most one clustered index per table
    - at most 999 non-clustered indexes per table
  - an index key (clustered / non-clustered)
    - at most 16 columns, 900 bytes\*

\* version differences

- clustered versus non-clustered index
  - clustered index
    - can be used for frequently executed queries
    - high degree of uniqueness
    - can be used in range queries
    - columns that are part of the search key:
      - shouldn't be frequently changed
      - should be narrow

- clustered versus non-clustered index
  - clustered index
    - organized as a B+ tree





- clustered versus non-clustered index
    - when creating a primary key on a table:
      - if a clustered index is not defined
      - a non-clustered index is not specified
- => a unique clustered index is created on the fields of the primary key

- unique indexes
  - such an index guarantees that the search key contains no duplicate values
  - specifying a unique index makes sense only when there are no entries with identical values in the key columns
  - uniqueness – useful information for the query optimizer

- key versus non-key index columns
  - key columns
    - the columns of the search key
  - non-key columns
    - columns specified in the INCLUDE clause when creating a non-clustered index

```
CREATE INDEX Index_Name  
    ON Schema_Name.Table_Name(Column)  
    INCLUDE (ColumnA, ColumnB, ColumnC)
```

- covering index
  - contains all the columns that are necessary in a query

- key versus non-key index columns
  - non-key columns - benefits
    - columns can be accessed from the index
    - data types that are not allowed in key columns can be used in non-key columns
    - non-key columns are not taken into account when computing the size of the key

- filtered indexes
  - optimized non-clustered indexes
  - can be used by queries that select from a certain subset of data

```
CREATE NONCLUSTERED INDEX IDX_eid_pid_f_od  
    ON Orders (EmpId, ProdId)  
    WHERE OrderDate IS NOT NULL
```

- benefits
  - better query performance
  - reduced index:
    - maintenance cost
    - storage cost

- index design
  - analyze the characteristics of the:
    - database
      - Online Transaction Processing (OLTP)
      - Online Analytical Processing (OLAP)
  - most frequently executed queries
  - columns used in queries

- index design - guidelines
  - database-related aspects
    - the presence of many indexes on a table deteriorates the performance of INSERT, UPDATE, DELETE, MERGE statements
    - indexing small tables is often useless

- index design - guidelines
  - query-related aspects
    - non-clustered indexes should be created on columns that are often used in WHERE and JOIN
    - covering indexes can significantly improve the performance of queries
    - as many records as possible should be changed in a single statement



- index design - guidelines
  - column-related aspects
    - length of the index key - as short as possible for clustered indexes
    - clustered indexes - better on unique / non-null columns
    - types ntext, text, image, varchar(max), nvarchar(max), varbinary(max) cannot be used for search key fields
    - column uniqueness
    - data distribution in the column
      - avoid indexes on columns with a small number of distinct values
    - consider creating filtered indexes for columns with well-defined subsets
    - order of columns in multicolumn indexes
      - first positions - columns in equality (=), inequality (>, <, BETWEEN) conditions
      - the rest of the columns should be ordered by level of distinctness

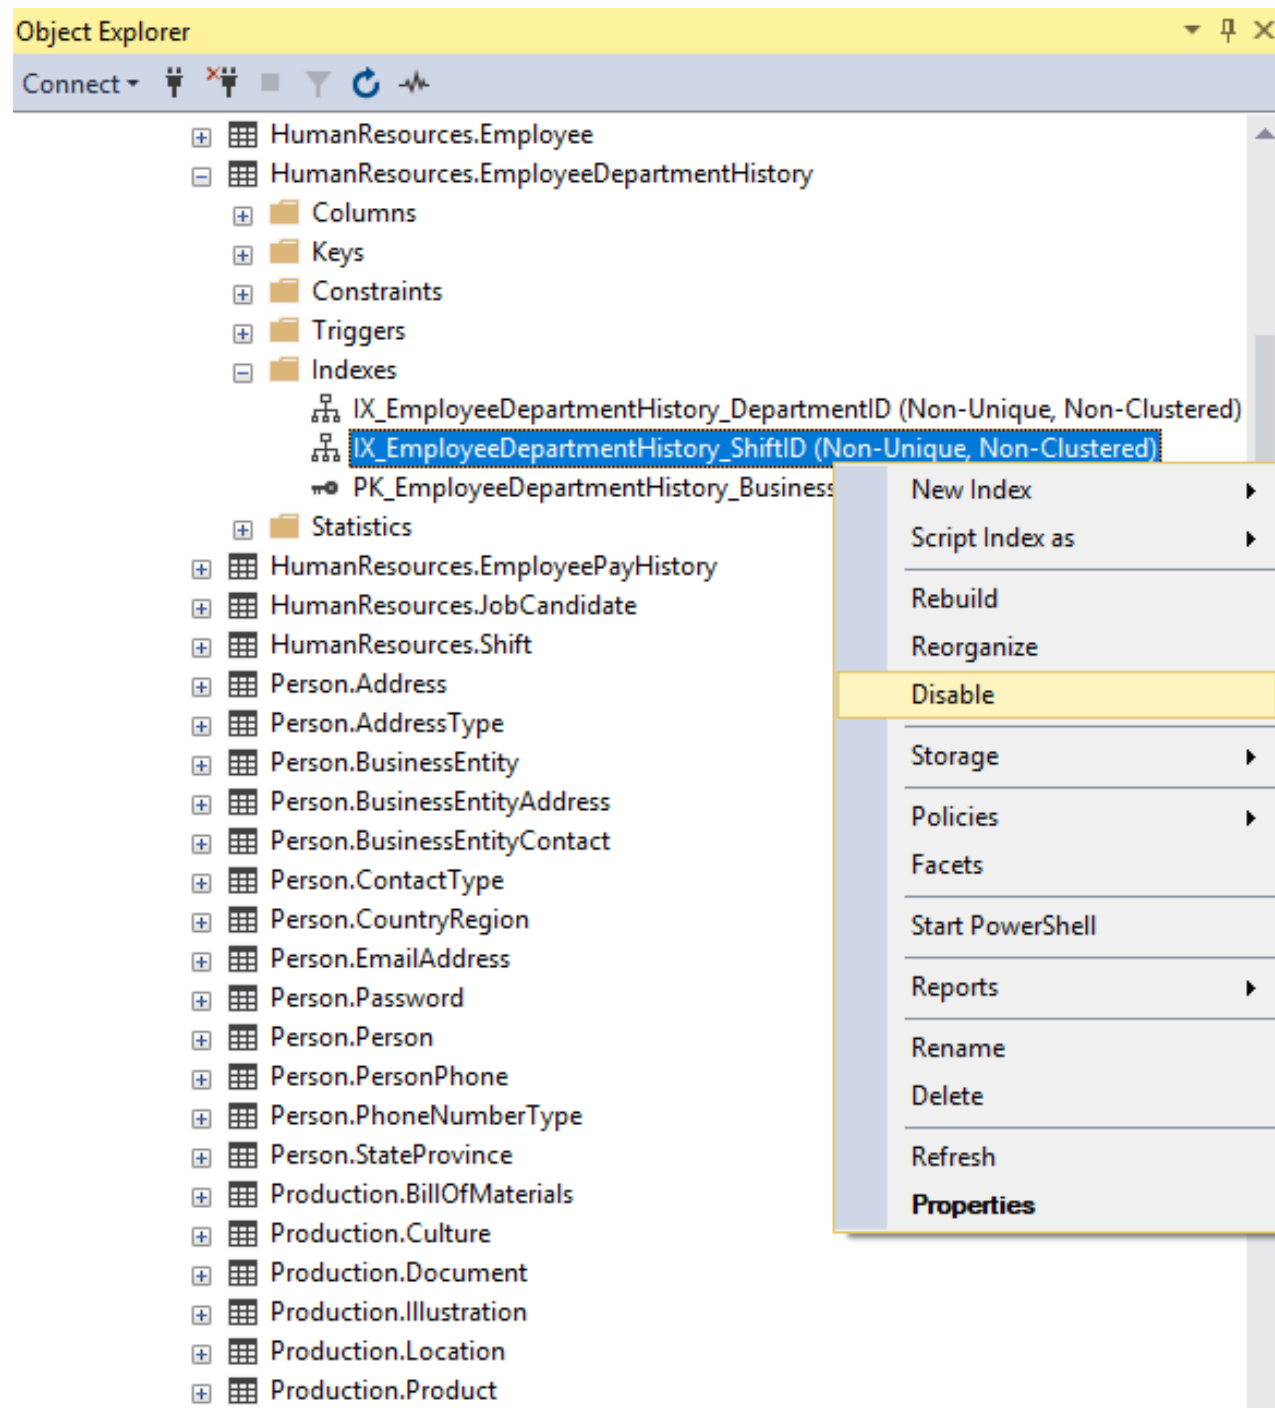
- index design - guidelines
  - column-related aspects
    - consider indexing computed columns

- indexes for deletes
  - when executing a delete statement
    - SQL Server searches for dependent rows by examining all the foreign keys, i.e., when a record *r* is deleted, the system checks whether *r* is referenced by other records
      - if there is an index, SQL Server uses it to check the existence of related data
      - if there isn't an index, the system has to perform a table scan
  - the performance of delete operations can be improved by creating indexes on foreign keys

- **disable indexes**

```
ALTER INDEX IX_EmployeeDepartmentHistory_ShiftID  
ON HumanResources.EmployeeDepartmentHistory DISABLE
```

- disable indexes



- enable indexes

```
ALTER INDEX IX_EmployeeDepartmentHistory_ShiftID  
    ON HumanResources.EmployeeDepartmentHistory REBUILD
```