

Database Management Systems

Lecture 7

Evaluating Relational Operators

Query Optimization (II)

- running example - schema
 - Students (SID: integer, SName: string, Age: integer)
 - Courses (CID: integer, CName: string, Description: string)
 - Exams (SID: integer, CID: integer, EDate: date, Grade: integer, FacultyMember: string)
- Students
 - every record has 50 bytes
 - there are 80 records / page
 - 500 pages of Students tuples
- Courses
 - every record has 50 bytes
 - there are 80 records / page
 - 100 pages of Courses tuples

- running example - schema
 - Students (SID: integer, SName: string, Age: integer)
 - Courses (CID: integer, CName: string, Description: string)
 - Exams (SID: integer, CID: integer, EDate: date, Grade: integer, FacultyMember: string)
- Exams
 - every record has 40 bytes
 - there are 100 records / page
 - 1000 pages of Exams tuples

* sorting

- can be explicitly required:
 - SELECT ... ORDER BY list
 - SELECT DISTINCT ...
- used by operators like:
 - join
 - union
 - intersection
 - set-difference
 - grouping


* sorting

e.g., the user wants to sort the collection of Courses records by name

- if the data to be sorted fits into available main memory:
 - use an internal sorting algorithm (Quick Sort or any other in-memory sorting algorithm can be used to sort a collection of records that fits into main memory)
- if the data to be sorted doesn't fit into available main memory:
 - use an external sorting algorithm
 - such an algorithm:
 - minimizes the cost of accessing the disk
 - breaks the data collection into subcollections of records
 - sorts the subcollections; a sorted subcollection of records is called a *run*
 - writes runs to disk (into temporary files)
 - merges runs

Simple Two-Way Merge Sort

- uses 3 buffer pages
- passes over the data multiple times
- can sort large data collections using a small amount of main memory



Simple Two-Way Merge Sort

- pass 0:
 - for each page P in the data collection:
 - read in page P -> sort page P -> save page P to disk

=> 1-page runs (runs that are 1 page long)

example: - read in the 1st page from Courses, sort the 80 records on it by course name, write out the sorted page to disk (i.e., a *run* that is one page long);

- read in the 2nd page from Courses, sort the 80 records on it by course name, write out sorted page to disk;


...

- read in the 100th page from Courses, sort the 80 records on it by course name, write out sorted page to disk

=> 100 1-page runs saved on disk

Simple Two-Way Merge Sort

- passes 1, 2, ... etc:
 - use 3 buffer pages
 - read and merge pairs of runs from the previous passes
 - produce runs that are twice as long



Simple Two-Way Merge Sort

- e.g., pass 1 (remember, pass 0 produced 100 1-page runs):
 - read in 2 runs from pass 0 (i.e., two pages holding Courses records, each of them sorted in pass 0), using 2 buffer pages
 - merge these runs writing to the 3rd available buffer page (the *output* buffer); when the output buffer fills up, write it out to disk (i.e., write a page of 80 sorted records to disk)
=> a run that is 2 pages long (it contains 160 Courses records, sorted by name)
 - read in and merge the next 2 runs from pass 0 ... => another run that is 2 pages long
 - continue while there are runs to be processed (read in and merged) from pass 0
 - at the end of pass 1 there are 50 2-page runs (each run consists of 2 pages holding 160 records sorted by course name)

Simple Two-Way Merge Sort

- another example – sort data collection (file of records) with 7 pages:

3, 4	6, 2	9, 4	8, 7	5, 6	3, 1	2
page 1	page 2	...				

- only the value of the key is displayed (the key on which the user wants to sort the collection, an integer number in the example)
- simplifying assumption that allows us to focus on the idea of the algorithm: a page can hold 2 records

- pass 0

- read in the collection one page at a time
- sort each page that is read in
- write out each sorted page to disk

=> 7 sorted runs that are 1 page long:

3, 4	2, 6	4, 9	7, 8	5, 6	1, 3	2
------	------	------	------	------	------	---

Simple Two-Way Merge Sort

- runs at the end of pass 0:

3, 4

2, 6

4, 9

7, 8

5, 6

1, 3

2

- pass 1

- read in & merge pairs of runs from pass 0
- produce runs that are twice as long
- read in runs 3, 4 and 2, 6 :
 - merge the runs and write to the output buffer
 - write the output buffer to disk one page at a time

=> run

2, 3

4, 6

- read in runs 4, 9 and 7, 8 :

- merge the runs and write to the output buffer
- write the output buffer to disk one page at a time

=> run

4, 7

8, 9

Simple Two-Way Merge Sort

- runs at the end of pass 0:

3, 4

2, 6

4, 9

7, 8

5, 6

1, 3

2

- pass 1

- read in runs 5, 6 and 1, 3 ...

=> run 1, 3 5, 6

- read in run 2 (the last run from pass 0) ...

=> run 2

=> 4 sorted runs that are 2 pages long (except for the last run):

2, 3 4, 6

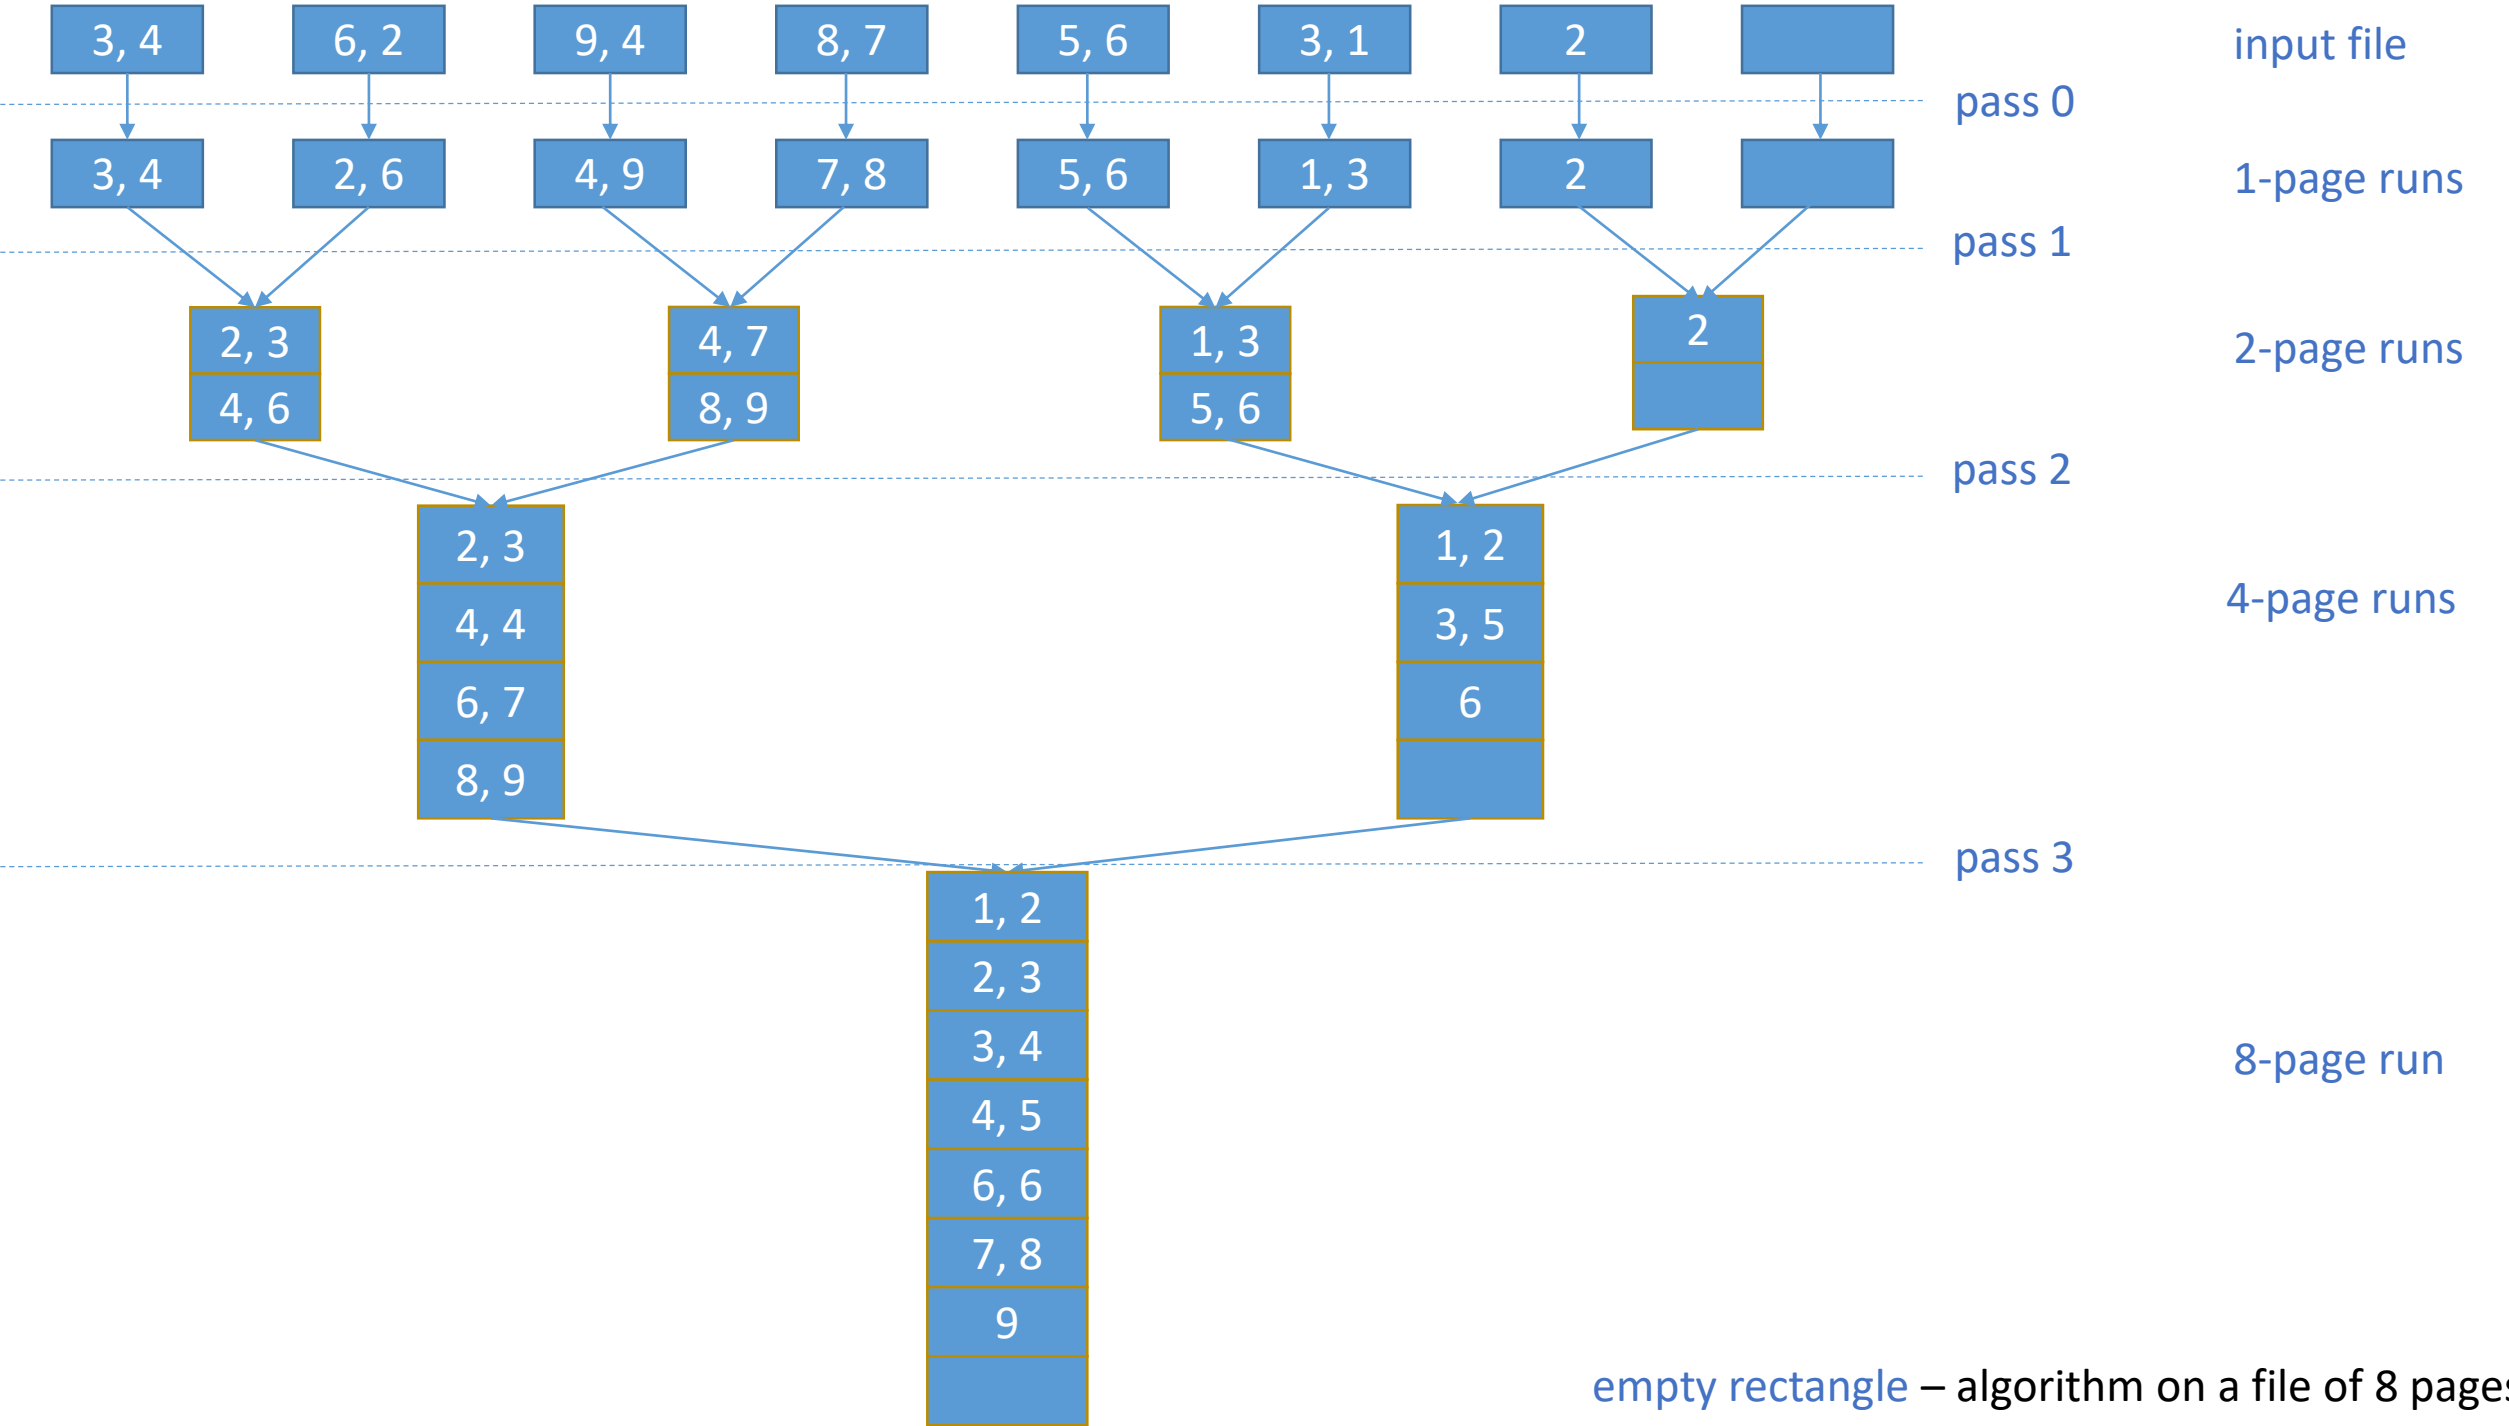
4, 7 8, 9

1, 3 5, 6

2

Simple Two-Way Merge Sort

- pass 2
 - read in & merge pairs of runs from pass 1
 - produce runs that are twice as long
- ...
- complete example, with all passes of the algorithm, on the next page ->



input file: 2^k pages

pass 0

=> 2^k sorted runs (1-page)

pass 1

=> 2^{k-1} sorted runs (2-pages)

pass 2

=> 2^{k-2} sorted runs (4-pages)


pass 3

=> 2^{k-3} sorted runs (8-pages)

i.e.,

pass k

=> one sorted run (2^k pages)



Simple Two-Way Merge Sort

- at each pass, each page in the input file is: read in, processed, and written out; there are 2 I/O operations per page, per pass


- number of passes:


- $\lceil \log_2 N \rceil + 1$

where N is the number of pages in the file to be sorted

- total cost:

- $2 * \text{number of pages} * \text{number of passes}$
 - $2 * N * (\lceil \log_2 N \rceil + 1)$ I/Os





- in each pass: read / process / write each page in the file
- number of passes:
 - $\lceil \log_2 N \rceil + 1$
- total cost:
 - $2 * N * (\lceil \log_2 N \rceil + 1)$ I/Os

- there are $N = 8$ pages, 4 passes
 - $\Rightarrow 2 * 8 * 4 = 64$ I/Os
 - or: $2 * 8 * (\lceil \log_2 8 \rceil + 1) = 2 * 8 * 4 = 64$ I/Os
- $N = 7$ pages, 4 passes
 - $\Rightarrow 2 * 7 * 4 = 56$ I/Os
 - or: $2 * 7 * (\lceil \log_2 7 \rceil + 1) = 56$ I/Os

External Merge Sort

- Simple Two-Way Merge Sort: buffer pages are not used effectively
 - for instance, if 200 buffer pages are available, this algorithm still uses only 2 input buffers for passes 1, 2, ...
- generalize the Two-Way Merge Sort algorithm to effectively use the available main memory and minimize the number of passes

->

External Merge Sort

- input file to be sorted: N pages
- B buffer pages are available
- pass 0:
 - use B buffer pages
 - read in B pages at a time and sort them in memory

=> $\left\lceil \frac{N}{B} \right\rceil$ runs of B pages each (except for the last one, which may be smaller)

->

External Merge Sort

- consider again the input file in the previous example:

3, 4	6, 2	9, 4	8, 7	5, 6	3, 1	2
page 1	page 2	...				

- $N = 7$ (number of pages in the file)
- $B = 4$ (there are 4 available buffer pages)

- pass 0 produces $\left\lceil \frac{N}{B} \right\rceil = \left\lceil \frac{7}{4} \right\rceil = 2$ runs:

- use all 4 buffer pages

- read in 4 pages:

3, 4	6, 2	9, 4	8, 7
------	------	------	------

- sort the pages in memory, write to disk a run that is 4 pages long:

2, 3	4, 4	6, 7	8, 9
------	------	------	------

- read in remaining 3 pages:


5, 6	3, 1	2
------	------	---

- sort pages in memory, write to disk run of 3 pages:

1, 2	3, 5	6
------	------	---

External Merge Sort

- input file to be sorted: N pages
- B buffer pages are available
- pass 1, 2 ...:
 - use $B-1$ pages for input, and one page for output
 - perform a $(B-1)$ -way merge in each pass (i.e., merge $B-1$ runs from the previous pass)



External Merge Sort

- runs at the end of pass 0:

2, 3	4, 4	6, 7	8, 9
------	------	------	------

1, 2	3, 5	6
------	------	---

- pass 1

- read in & merge the first $B-1 = 4-1 = 3$ runs from pass 0
- pass 0 produced only 2 runs in this example; read in and merge these 2 runs:

=> run

1, 2	2, 3	3, 4	4, 5	6, 6	7, 8	9
------	------	------	------	------	------	---

External Merge Sort

- another example:
 - 5 buffer pages $B = 5$
 - sort file with 108 pages $N = 108$
- pass 0
 - use all 5 buffer pages
 - read in the first 5 pages of the file, sort them in memory, write the resulting run to disk (5 pages long)
 - read in the next 5 pages of the file, sort them in memory, write the resulting run to disk (5 pages long)
 - ...
 - read in the remaining 3 pages of the file, sort them in memory, write the resulting run to disk (3 pages long)
 - 21 runs are 5 pages long; 1 run is 3 pages long

External Merge Sort

- another example: $B = 5$, $N = 108$
- pass 0
 - at the end of pass 0 there are $\left\lceil \frac{N}{B} \right\rceil = \left\lceil \frac{108}{5} \right\rceil = 22$ runs
- pass 1
 - use $B-1 = 5-1 = 4$ pages for input, and one page for output
 - do a 4-way merge: read in and merge 4 runs from the previous pass
 - read in the first 4 runs from pass 0 (each run into an input buffer)
 - merge the runs and write to the output buffer
 - write the output buffer to disk one page at a time
 - => a run that is 20 pages long (4 runs from pass 0 times 5 pages per run)
 - read in the next 4 runs from pass 0; merge the runs and write to the output buffer; write the output buffer to disk one page at a time
 - => another run (20 pages long)

...

External Merge Sort

- another example: $B = 5$, $N = 108$
- pass 0
 - at the end of pass 0 there are 22 runs
- pass 1
 - read in the last 2 runs from pass 0 (one has 5 pages, the other one has 3 pages)
 - merge the runs and write to the output buffer; write the output buffer to disk one page at a time
 - => the last run (8 pages long)
 - at the end of pass 1 there are $\left\lceil \frac{22}{4} \right\rceil = 6$ runs
 - 5 runs are 20 pages long; 1 run is 8 pages long

External Merge Sort

- another example: $B = 5$, $N = 108$
- pass 1
 - at the end of pass 1 there are 6 runs
- pass 2
 - 4-way merge
 - read in the first 4 runs from pass 1
 - merge the runs and write to the output buffer; write the output buffer to disk one page at a time
 - => a run that is 80 pages long (4 runs from pass 1 times 20 pages per run)
 - read in the remaining 2 runs from pass 1 (20 and 8 pages, respectively)
 - => a run that is 28 pages long
 - at the end of pass 2 there are $\left\lceil \frac{6}{4} \right\rceil = 2$ runs

External Merge Sort

- another example: $B = 5$, $N = 108$
- pass 2
 - at the end of pass 2 there are 2 runs
- pass 3
 - read in the 2 runs from pass 2 and merge them
=> a run that is 108 pages long, representing the sorted file

External Merge Sort

- cost
 - N – number of pages in the input file, B – number of available pages in the buffer
 - at each pass: read / process / write each page
 - number of passes: $\lceil \log_{B-1} \lceil N/B \rceil \rceil + 1$
 - total cost: $2 * N * \left(\lceil \log_{B-1} \left\lceil \frac{N}{B} \right\rceil \rceil + 1 \right)$ I/Os
- previous example: B = 5 and N = 108, with 4 passes over the data
 - cost:
$$2 * 108 * 4 = 864 \text{ I/Os}$$
 - $2 * 108 * \left(\lceil \log_{5-1} \left\lceil \frac{108}{5} \right\rceil \rceil + 1 \right) = 216 * (\lceil \log_4 22 \rceil + 1) = 216 * 4 = 864 \text{ I/Os}$

- B buffer pages
- sort file with N pages

Simple Two-Way Merge Sort

pass 0 => N runs

number of passes = $\lceil \log_2 N \rceil + 1$

External Merge Sort

pass 0 => $\left\lceil \frac{N}{B} \right\rceil$ runs


number of passes = $\left\lceil \log_{B-1} \left\lceil \frac{N}{B} \right\rceil \right\rceil + 1$

- External Merge Sort – reduced number of:
 - runs produces by the 1st pass
 - passes over the data
- B is usually large => significant performance gains

External Merge Sort – number of passes for different values of N and B

N	B = 3	B = 5	B = 9	B = 17	B = 129	B = 257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

- * minimize the number of runs - optional
- external merge sort
 - N pages in the file, B buffer pages $\Rightarrow \lceil N/B \rceil$ runs of B pages each
- improvement
 - algorithm known to produce sorted runs of approximately $2*B$ pages (on average)
 - use 1 page as an input buffer, 1 page as an output buffer
 - the remaining buffer pages are collectively referred to as the *current set*
- example - sort file in ascending order on some key k :
- repeatedly pick record r with smallest k in current set such that k is $>$ the largest k in the output buffer
- append r to output buffer (the output buffer is kept sorted)
- use the extra space in the current set to bring in the next tuple from the input buffer



* minimize the number of runs - optional

- process all tuples in the input buffer, then read in the next page of the file
- when the output buffer fills up, write it to disk (add its content to the run that is currently being built)
- the current run is completed when every k value in the current set is $<$ the largest k value in the output buffer; when this happens, the output buffer is written out (its content becomes the last page in the current run), and a new run is started

Sort-Merge Join

- equality join, one join column: $E \bowtie_{i=j} S$ (i^{th} column's value in $E = j^{\text{th}}$ column's value in S)
- sort E and S on the join column (if not already sorted):
 - for instance, by using External Merge Sort

\Rightarrow *partitions* = groups of tuples with the same value in the join column
- merge E and S ; look for tuples e in E , s in S such that $e_i = s_j$:
 - while *current* $e_i <$ *current* s_j
 - advance the scan of E
 - while *current* $e_i >$ *current* s_j
 - advance the scan of S
 - if *current* $e_i =$ *current* s_j
 - output joined tuples $\langle e, s \rangle$, where e and s are in the current partition (i.e., they have the same value in the i^{th} and j^{th} column, respectively)
 - there could be multiple tuples in E with the same value in the i^{th} column as the current tuple e (same is true for S)

Sort-Merge Join

- partitions are illustrated on tables Students and Exams below (join column SID in both tables):

SID	SName	Age
20	Ana	20
30	Dana	20
40	Dan	20
45	Daniel	20
50	Ina	20

SID	CID	EDate	Grade	FacultyMember
30	2	20/1/2018	10	Ionescu
30	1	21/1/2018	9.99	Pop
45	2	20/1/2018	9.98	Ionescu
45	1	21/1/2018	9.98	Pop
45	3	22/1/2018	10	Stan
50	2	20/1/2018	10	Ionescu

Sort-Merge Join

- during the merging phase, E is scanned once; every partition in S is scanned as many times as there are matching tuples in the corresponding partition in E

E

...	i th column	...
	1	
	2	
	2	
	2	
	3	
	...	

S

...	j th column	...
	0	
	2	partition P
	2	
	4	
	...	
	...	

- for instance, partition P in the above table S is scanned 3 times, once per matching tuple in the corresponding partition in E
- there are 6 output joined tuples $\langle e, s \rangle$ for partition P
- this algorithm avoids the enumeration of the cross-product: tuples in a partition in E are compared only with the S tuples in the same partition!

Sort-Merge Join

- cost:
 - sorting E
 - cost: $O(M \log M)$
 - sorting S
 - cost: $O(N \log N)$
 - cost of merging: $M + N$ I/Os, assuming partitions in S are scanned only once
 - worst-case scenario: $M * N$ I/Os (when all records in E and S have the same value in the join column)

* E - M pages; S - N pages*

Sort-Merge Join ($\text{Exams} \bowtie_{\text{Exams.SID}=\text{Students.SID}} \text{Students}$)

- 100 buffer pages
 - sort Exams
 - 2 passes \Rightarrow cost: $2 * 2 * 1000 = 4000$ I/Os
 - sort Students
 - 2 passes \Rightarrow cost: $2 * 2 * 500 = 2000$ I/Os
 - merging phase
 - cost: $1000 + 500 = 1500$ I/Os
 - total cost: $4000 + 2000 + 1500 = 7500$ I/Os
 - similar to the cost of Block Nested Loops Join

* E - M pages, p_E records / page *

* 1000 pages * * 100 records / page *

* S - N pages, p_S records / page *

* 500 pages * * 80 records / page *

Sort-Merge Join ($\text{Exams} \bowtie_{\text{Exams.SID}=\text{Students.SID}} \text{Students}$)

- 35 buffer pages
 - sort Exams
 - 2 passes \Rightarrow cost: $2 * 2 * 1000 = 4000$ I/Os
 - sort Students
 - 2 passes \Rightarrow cost: $2 * 2 * 500 = 2000$ I/Os
 - merging phase
 - cost: $1000 + 500 = 1500$ I/Os
 - total cost: $4000 + 2000 + 1500 = 7500$ I/Os
 - ex: compute cost of BNLJ and compare

* E - M pages, p_E records / page *

* 1000 pages * * 100 records / page *

* S - N pages, p_S records / page *

* 500 pages * * 80 records / page *

Sort-Merge Join ($\text{Exams} \bowtie_{\text{Exams.SID}=\text{Students.SID}} \text{Students}$)

- 300 buffer pages
 - sort Exams
 - 2 passes \Rightarrow cost: $2 * 2 * 1000 = 4000$ I/Os
 - sort Students
 - 2 passes \Rightarrow cost: $2 * 2 * 500 = 2000$ I/Os
 - merging phase
 - cost: $1000 + 500 = 1500$ I/Os
 - total cost: $4000 + 2000 + 1500 = 7500$ I/Os
 - ex: compute cost of BNLJ and compare

* E - M pages, p_E records / page *

* 1000 pages * * 100 records / page *

* S - N pages, p_S records / page *

* 500 pages * * 80 records / page *

References

- [Ra00] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (2nd Edition), McGraw-Hill, 2000
- [Da03] DATE, C.J., An Introduction to Database Systems (8th Edition), Addison-Wesley, 2003
- [Ga08] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book, Prentice Hall Press, 2008
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007, <http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si10] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, McGraw-Hill, 2010, <http://codex.cs.yale.edu/avi/db-book/>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems, <http://infolab.stanford.edu/~ullman/fcdb.html>