

# Database Management Systems

Lecture 8

Evaluating Relational Operators

Query Optimization (III)

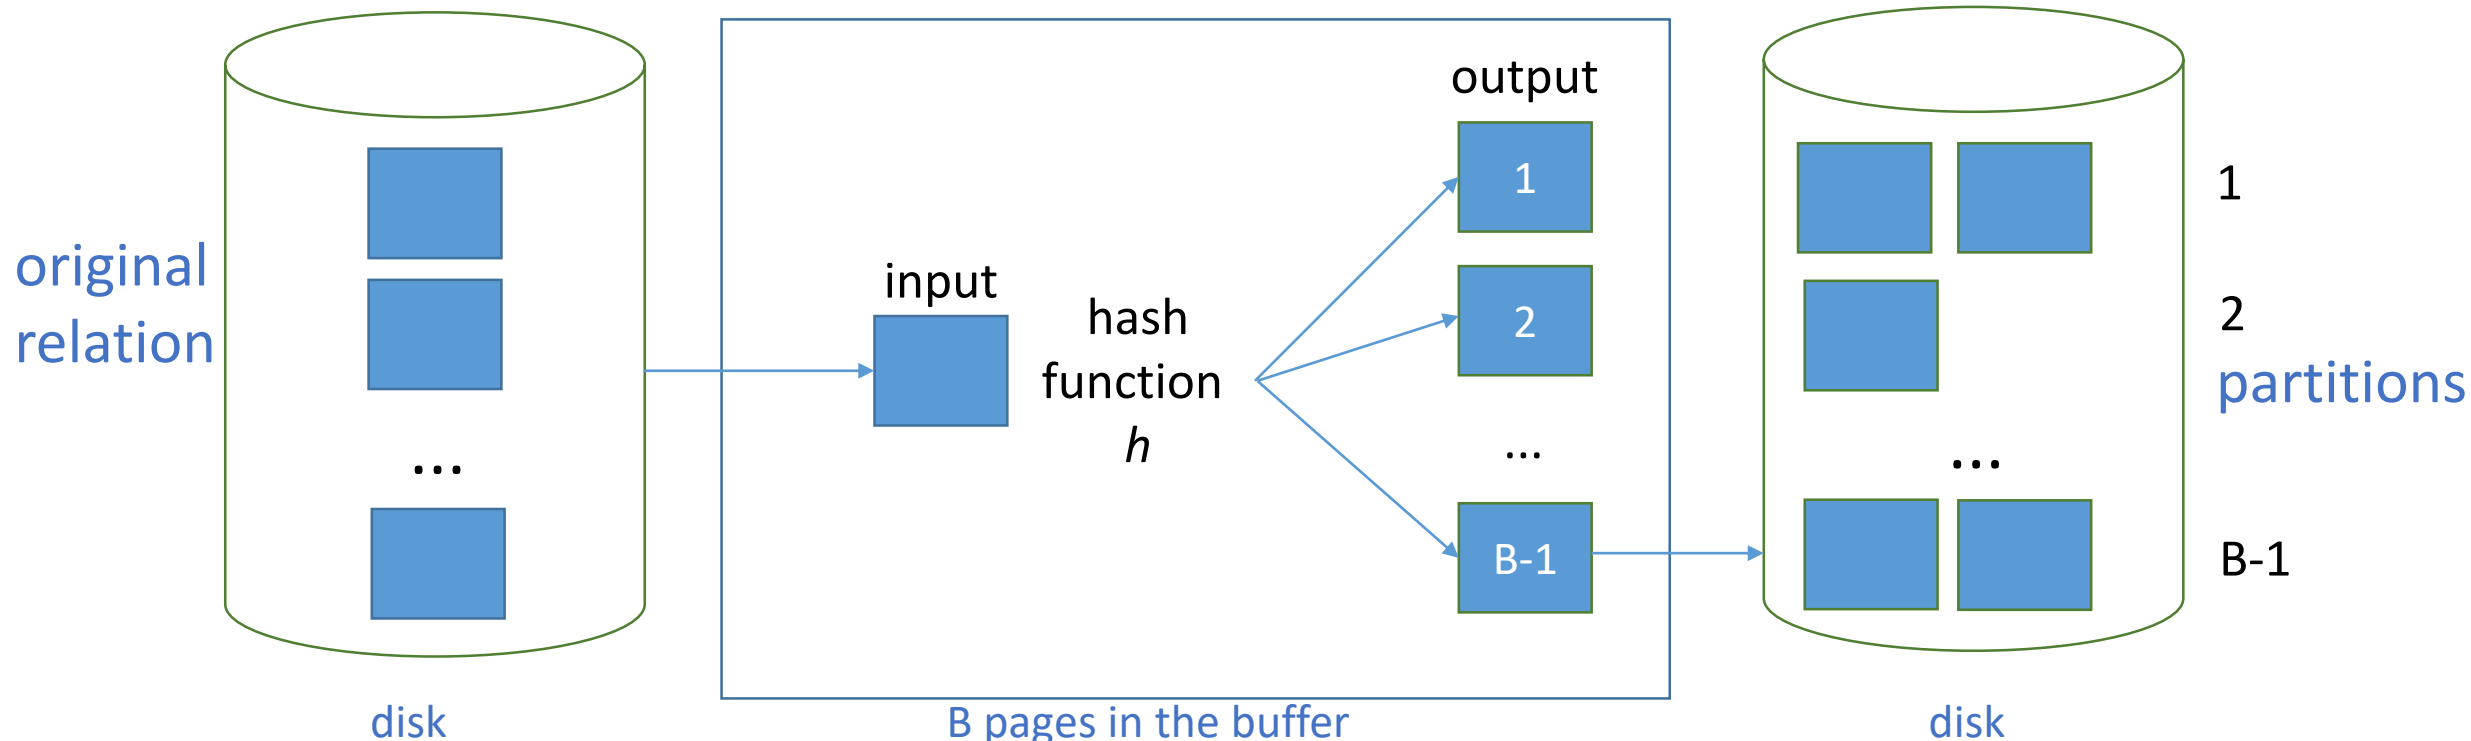
- running example - schema
  - Students (SID: integer, SName: string, Age: integer)
  - Courses (CID: integer, CName: string, Description: string)
  - Exams (SID: integer, CID: integer, EDate: date, Grade: integer, FacultyMember: string)
- Students
  - every record has 50 bytes
  - there are 80 records / page
  - 500 pages of Students tuples
- Courses
  - every record has 50 bytes
  - there are 80 records / page
  - 100 pages of Courses tuples
- Exams
  - every record has 40 bytes
  - there are 100 records / page
  - 1000 pages of Exams tuples

Hash Join - equality join, one join column:  $E \bowtie_{i=j} S$

- phases: partitioning (building phase) & probing (matching phase)
- partitioning phase:
  - there are  $B$  pages available in the buffer:
    - use one page as the input buffer page
    - and the remaining  $B-1$  pages as output buffer pages
  - choose a hash function  $h$  that distributes tuples uniformly to one of  $B-1$  partitions
  - hash  $E$  and  $S$  on the join column (the  $i^{\text{th}}$  column of  $E$ , the  $j^{\text{th}}$  column of  $S$ ) with the same hash function  $h$

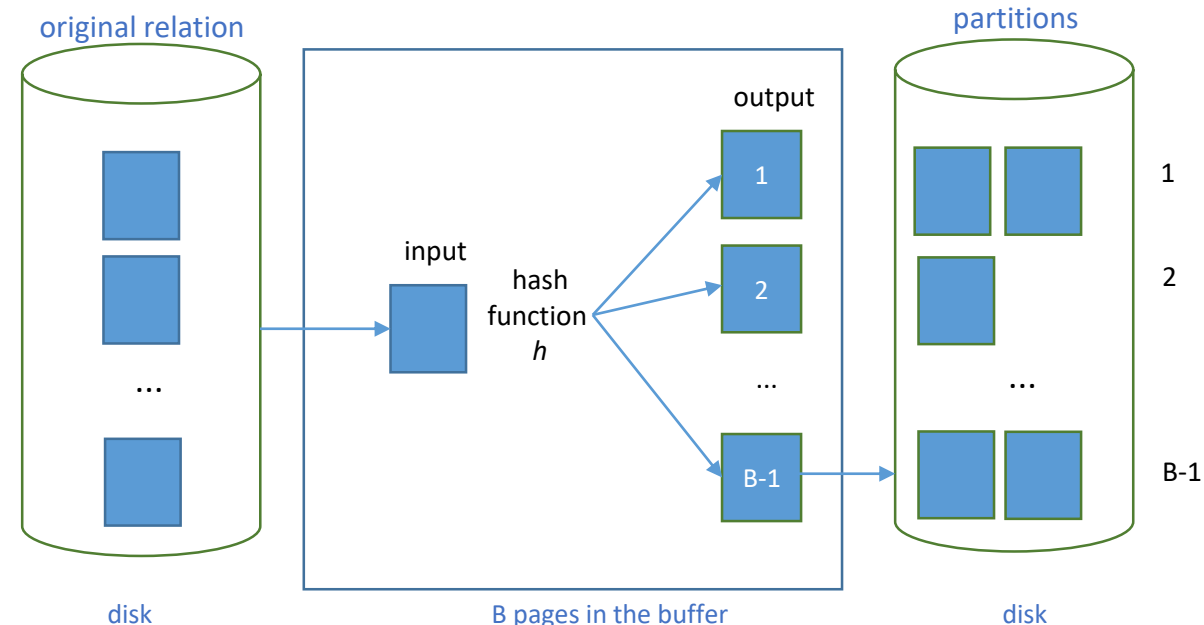
## Hash Join

- hash E on the join column with hash function  $h$  (similarly for S):
  - for each tuple  $e$  in E, compute  $h(e_i)$   
( $e_i$ : the value of the  $i^{\text{th}}$  column in tuple  $e$ )
  - add tuple  $e$  to the output buffer page that it is hashed to by  $h$  (buffer page  $h(e_i)$ )
  - when an output buffer page fills up, flush the page to disk



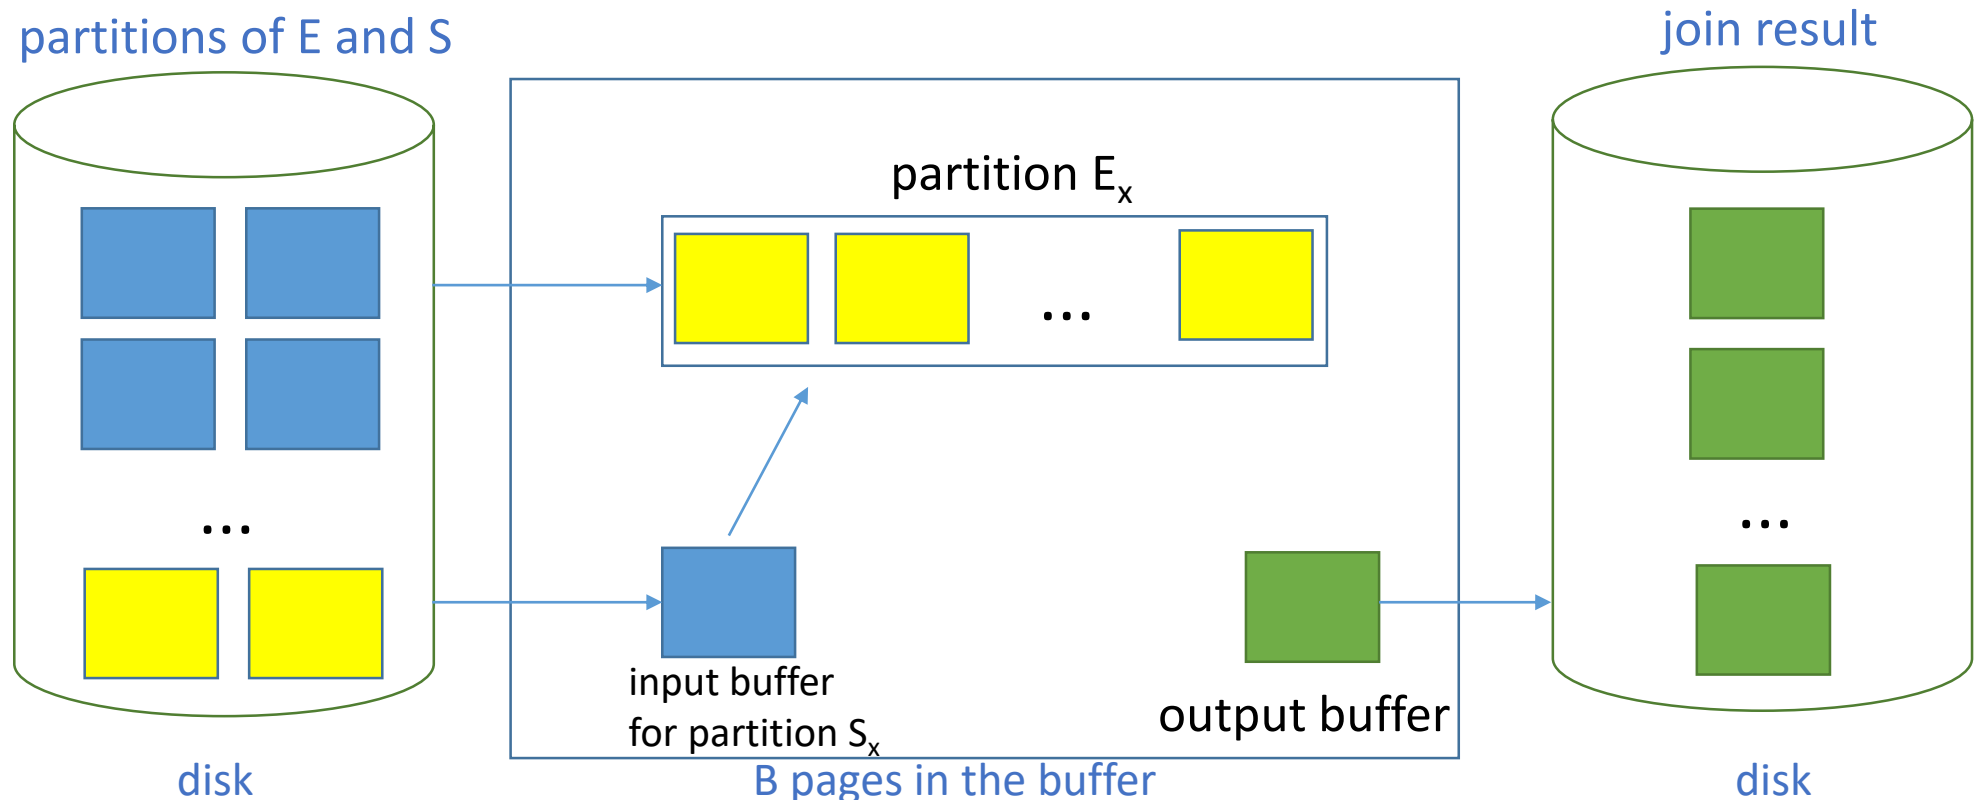
## Hash Join

- partitioning phase  $\Rightarrow$  *partitions* of  $E$  ( $E_1, E_2$ , etc) and  $S$  ( $S_1, S_2$ , etc) on disk
- partition = collection of tuples that have the same hash value
- tuples in partition  $E_1$  can only join with tuples in partition  $S_1$  (they cannot join with tuples in partitions  $S_2$  or  $S_3$ , for instance, since these tuples have a different hash value)
- so to compute the join, we need to scan  $E$  and  $S$  only once (provided any partition of  $E$  fits in main memory)
- when reading in a partition  $E_k$  of  $E$ , we must scan only the corresponding partition  $S_k$  of  $S$  to find matching tuples (compare tuples  $e$  in  $E_k$  with tuples  $s$  in  $S_k$  to test the join condition *value of  $i^{\text{th}}$  column in  $E$  = value of  $j^{\text{th}}$  column in  $S$* )



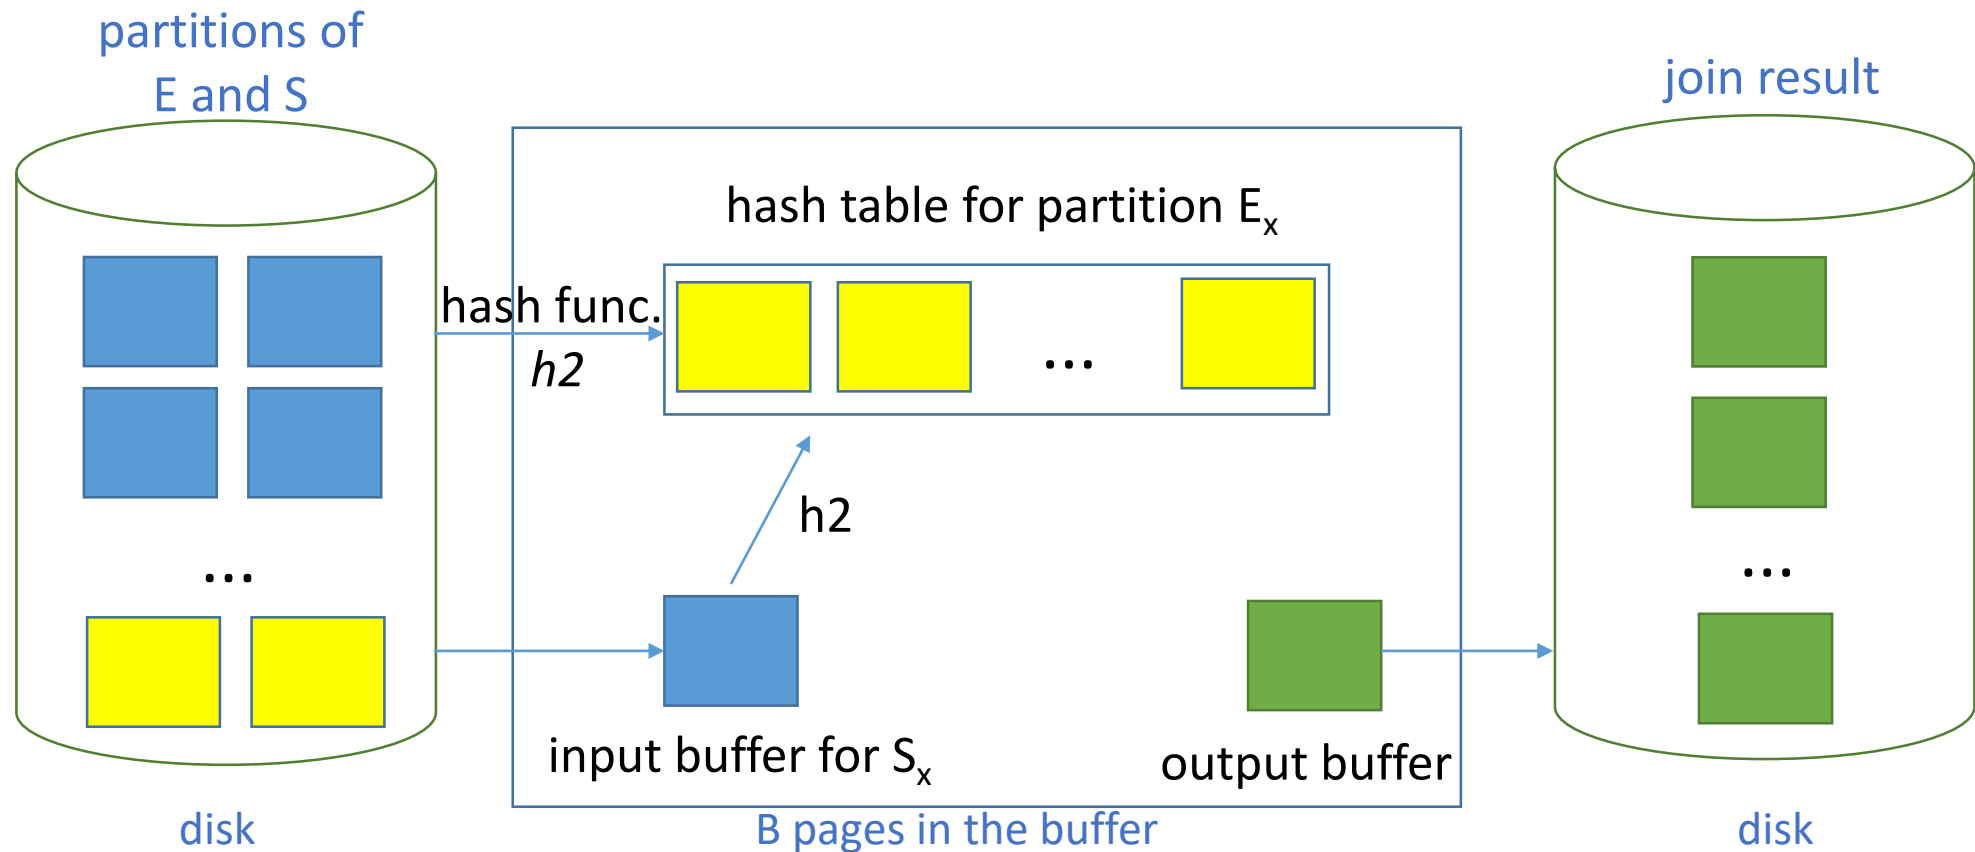
## Hash Join

- probing phase:
  - read in a partition of the smaller relation (e.g., E) and scan the corresponding partition of S for matching tuples
  - use one page as the input buffer for S, one page as the output buffer, and the remaining pages to read in partitions of E



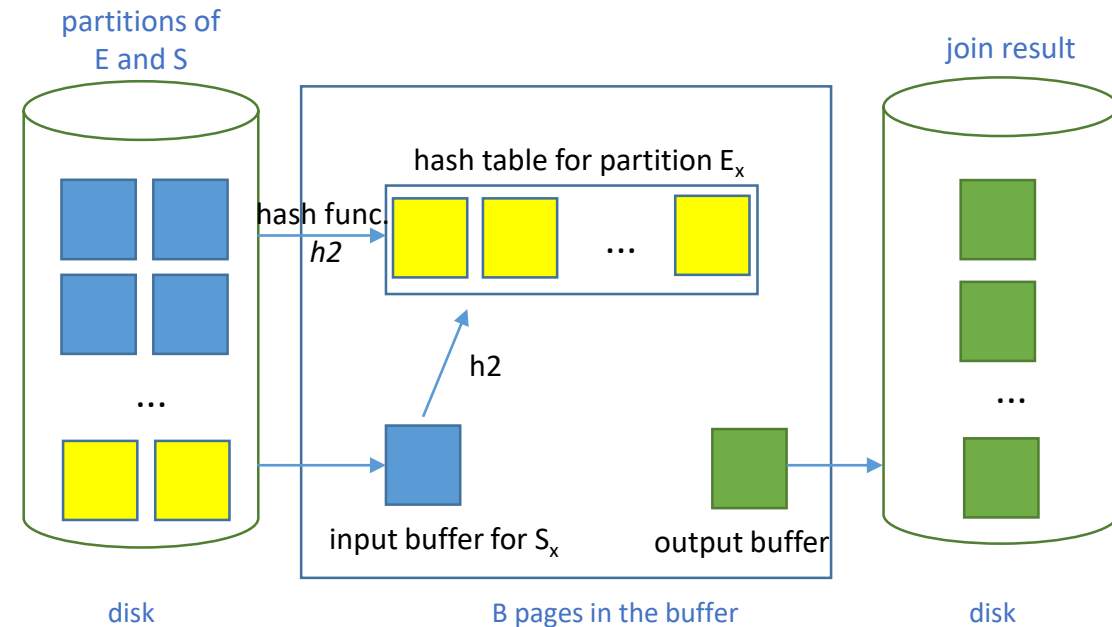
## Hash Join

- probing phase:
  - in practice, to reduce CPU costs, an in-memory hash table is built, using a different function  $h2$ , for the E partition



## Hash Join

- probing phase:
  - in practice, to reduce CPU costs, an in-memory hash table is built, using a different function  $h2$ , for the E partition
- consider a partition  $E_x$  of E
- build in-memory hash table for  $E_x$  using hash function  $h2$  (the function is applied to the join column of E)
- for each tuple  $s$  in partition  $S_x$ , find matching tuples in the hash table using the hash value  $h2(s_j)$
- result tuples  $\langle e, s \rangle$  are written to output buffer
- once partitions  $E_x$  and  $S_x$  are processed, the hash table is emptied (to prepare for the next partition)





## Hash Join

- cost:
    - partitioning:
      - both E and S are read and written once => cost:  $2*(M+N)$  I/Os
    - probing:
      - scan each partition once => cost:  $M+N$  I/Os
- => total cost:  $3*(M+N)$  I/Os
- assumption: each partition fits into memory during probing
  - $3*(1000 + 500) = 4500$  I/Os

\* E - M pages,  $p_E$  records / page \*

\* S - N pages,  $p_S$  records / page \*

\* 1000 pages \* \* 100 records / page \*

\* 500 pages \* \* 80 records / page \*

## Hash Join

- *partition overflow* – an E partition does not fit in memory during probing:  
apply hash join technique recursively:
  - divide E, S into subpartitions
  - join subpartitions pairwise
  - if subpartitions don't fit in memory, apply hash join technique recursively

## Hash Join

- memory requirements - objective: partition in E fits into main memory (S - similarly)
  - B buffer pages; need one input buffer  $\Rightarrow$  maximum number of partitions: B-1
  - size of largest partition: B - 2 (need one input buffer for S, one output buffer)
  - assume uniformly sized partitions  $\Rightarrow$  size of each E partition:  $M/(B-1)$   
 $\Rightarrow M/(B-1) < B-2 \Rightarrow$  we need approximately  $B > \sqrt{M}$
- if an in-memory hash table is used to speed up tuple matching  $\Rightarrow$  need a little more memory (because the hash table for a collection of tuples will be a little larger than the collection itself)

\* E - M pages,  $p_E$  records / page \*

\* 1000 pages \*

\* 100 records / page\*

## general join conditions

- equalities over several attributes
  - $E.SID = S.SID \text{ AND } E.attrE = S.attrS$ 
    - index nested loops join
      - Exams – inner relation:
        - build index on Exams with search key  $\langle SID, attrE \rangle$  (if not already created)
        - can also use index on SID or index on attrE
      - Students – inner relation (similar)
  - sort-merge join
    - sort Exams on  $\langle SID, attrE \rangle$ , sort Students on  $\langle SID, attrS \rangle$
  - hash join
    - partition Exams on  $\langle SID, attrE \rangle$ , partition Students on  $\langle SID, attrS \rangle$
  - other join algorithms
    - essentially unaffected

## general join conditions

- inequality comparison
  - $E.attrE < S.attrS$ 
    - index nested loops join
      - B+ tree index required
    - sort-merge join
      - not applicable
    - hash join
      - not applicable
    - other join algorithms
      - essentially unaffected

- \* no join algorithm is uniformly superior to others
- choice of a good algorithm depends on:
  - size(s) of:
    - joined relations
    - buffer pool
  - available access methods

## Selection

Q:

```
SELECT *  
FROM Exams E  
WHERE E.FacultyMember = 'Ionescu'
```

- use information in the selection condition to reduce the number of retrieved tuples
- e.g.,  $|Q| = 4$  (result set has 4 tuples), there's a B+ tree index on FacultyMember
  - it's expensive to scan E (1000 I/Os) to evaluate the query
  - should use the index instead
- selection algorithms based on the following techniques:
  - iteration, indexing

\* E - M pages,  $p_E$  records / page \*      \* 1000 pages \* \* 100 records / page\*

## Selection

- simple selections\*
  - $\sigma_{E.attr \text{ op } val}(E)$
- no index on *attr*, data not sorted on *attr*
  - must scan E and test the condition for each tuple
  - access path: file scan

=> cost: M I/Os = 1000 I/Os
- no index, sorted data (E physically sorted on *attr*)
  - binary search to locate 1<sup>st</sup> tuple that satisfies condition  
and
  - scan E starting at this position until condition is no longer satisfied
  - access method: sorted file scan



## Selection

- simple selections
  - $\sigma_{E.attr \text{ op } val}(E)$
- no index, sorted data (E physically sorted on *attr*)  
=> cost:
  - binary search:  $O(\log_2 M)$
  - scan cost: varies from 0 to M
  - binary search on E
    - $\log_2 1000 \approx 10$  I/Os

## Selection

- simple selections
  - $\sigma_{E.attr \text{ op } val}(E)$
- B+ tree index on *attr*
  - \* search tree to find 1<sup>st</sup> index\* entry pointing to a qualifying E tuple
    - cost: typically 2, 3 I/Os
  - \* scan leaf pages to retrieve all qualifying entries
    - cost: depends on the number of qualifying entries
  - \* for each qualifying entry - retrieve corresponding tuple in E
    - cost: depends on the number of tuples and the nature of the index (clustered / non-clustered)

\*Review *Indexes* - lecture notes (*Databases* course)

## Selection

- simple selections
  - $\sigma_{E.attr \text{ op } val}(E)$
- B+ tree index on *attr*
  - assumption
    - indexes use a2 or a3
    - a1-based index => data entry contains the data record => the cost of retrieving records = the cost of retrieving the data entries!
  - access path: B+ tree index
    - clustered index:
      - best access path when *op* is not *equality*
      - good access path when *op* is *equality*

## Selection

- simple selections:  $\sigma_{E.attr \text{ op } val}(E)$

- B+ tree index on *attr*

Q

```
SELECT *
```

```
FROM Exams E
```

```
WHERE E.FacultyMember < 'C%'
```

- names uniformly distributed with respect to 1<sup>st</sup> letter

=>  $|Q| \approx 10,000$  tuples = 100 pages

- clustered B+ tree index on FacultyMember

=> cost of retrieving tuples:  $\approx 100$  I/Os (a few I/Os to get from root to leaf)

- non-clustered B+ tree index on FacultyMember

=> cost of retrieving tuples: up to 1 I/O per tuple (worst case) => up to 10.000 I/Os

\* E - M pages,  $p_E$  records / page \*

\* 1000 pages \* \* 100 records / page\*

## Selection

- simple selections:  $\sigma_{E.attr \text{ op } val}(E)$

- B+ tree index on *attr*

```
SELECT *
```

```
FROM Exams E
```

```
WHERE E.FacultyMember < 'C%'
```

- non-clustered B+ tree index on FacultyMember
  - refinement - sort rids in qualifying data entries by page-id\*  
=> a page containing qualifying tuples is retrieved only once
  - cost of retrieving tuples: number of pages containing qualifying tuples (but such tuples are probably stored on more than 100 pages)
- range selections
  - non-clustered indexes can be expensive
  - could be less costly to scan the relation (in our example: 1000 I/Os)

\*Review DB – Physical Structure - lecture notes (Databases course)

## Selection

- general selections
  - selections without disjunctions
- C - CNF condition without disjunctions
  - evaluation options:
    1. use the most selective access path
      - if it's an index I:
        - apply conjuncts in C that match I
        - apply rest of conjuncts to retrieved tuples
      - example
        - $c < 100 \text{ AND } a = 3 \text{ AND } b = 5$ 
          - can use a B+ tree index on  $c$  and check  $a = 3 \text{ AND } b = 5$  for each retrieved tuple
          - can use a hash index on  $a$  and  $b$  and check  $c < 100$  for each retrieved tuple

## Selection

- general selections - selections without disjunctions
  - evaluation options:
    2. use several indexes - when several conjuncts match indexes using a2 / a3
      - compute sets of rids of candidate tuples using indexes
      - intersect sets of rids, retrieve corresponding tuples
      - apply remaining conjuncts (if any)
      - example:  $c < 100 \text{ AND } a = 3 \text{ AND } b = 5$ 
        - use a B+ tree index on  $c$  to obtain rids of records that meet condition  $c < 100$  ( $R_1$ )
        - use a hash index on  $a$  to retrieve rids of records that meet condition  $a = 3$  ( $R_2$ )
        - compute  $R_1 \cap R_2 = R_{int}$
        - retrieve records with rids in  $R_{int}$  ( $R$ )
        - check  $b = 5$  for each record in  $R$

## Selection

- general selections
  - selections with disjunctions
- C - CNF condition with disjunctions, i.e., some conjunct  $J$  is a disjunction of terms
  - if some term  $T$  in  $J$  requires a file scan, testing  $J$  by itself requires a file scan
    - example:  $a < 100 \vee b = 5$ 
      - hash index on  $b$ , hash index on  $c$
  - => check both terms using a file scan (i.e., best access path: file scan)
- compare with the example below:
  - $(a < 100 \vee b = 5) \wedge c = 7$
  - hash index on  $b$ , hash index on  $c$
  - => use index on  $c$ , apply  $a < 100 \vee b = 5$  to each retrieved tuple (i.e., most selective access path: index)



## Selection

- general selections
  - selections with disjunctions
- C - CNF condition with disjunctions
  - every term  $T$  in a disjunction matches an index

=> retrieve tuples using indexes, compute union

  - example
    - $a < 100 \vee b = 5$
    - B+ tree indexes on  $a$  and  $b$
    - use index on  $a$  to retrieve records that meet condition  $a < 100$  ( $R_1$ )
    - use index on  $b$  to retrieve records that meet condition  $b = 5$  ( $R_2$ )
    - compute  $R_1 \cup R_2 = R$
    - if all matching indexes use a2 or a3 => take union of rids, retrieve corresponding tuples

## Projection

- $\Pi_{\text{SID, CID}}(\text{Exams})$

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- to implement projection:
  - eliminate:
    - unwanted columns
    - duplicates
- projection algorithms - *partitioning* technique:
  - sorting
  - hashing

## Projection Based on Sorting

- step 1
  - scan  $E \Rightarrow$  set of tuples containing only desired attributes ( $E'$ )
  - cost:
    - scan  $E$ :  $M$  I/Os
    - write temporary relation  $E'$ :  $T$  I/Os
      - $T$  depends on: number of columns and their sizes,  $T$  is  $O(M)$
- step 2
  - sort tuples in  $E'$
  - sort key: all columns
  - cost:  $O(T \log T)$  (also  $O(M \log M)$ )
- step 3
  - scan sorted  $E'$ , compare adjacent tuples, eliminate duplicates
  - cost:  $T$
- total cost:  $O(M \log M)$

## Projection Based on Sorting

### \* example

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- scan Exams: 1000 I/Os
- size of tuple in E': 10 bytes

=> cost of writing temporary relation E': 250 I/Os

- available buffer pages: 20
  - E' can be sorted in 2 passes
  - sorting cost:  $2 * 2 * 250 = 1000$  I/Os
- final scan of E' - cost: 250 I/Os

=> total cost:  $1000 + 250 + 1000 + 250 = 2500$  I/Os

\* E – record size = 40 bytes \*                      \* 1000 pages \*                      \* 100 records / page\*

## Projection Based on Sorting

### \* example

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- scan Exams: 1000 I/Os
- size of tuple in E': 10 bytes

=> cost of writing temporary relation E': 250 I/Os

- available buffer pages: 257
  - E' can be sorted in 1 pass
  - sorting cost:  $2 * 1 * 250 = 500$  I/Os
- final scan of E' - cost: 250 I/Os

=> total cost:  $1000 + 250 + 500 + 250 = 2000$  I/Os

\* E – record size = 40 bytes \*                      \* 1000 pages \*                      \* 100 records / page\*

## Projection Based on Sorting

- improvement
  - adapt the sorting algorithm to do projection with duplicate elimination
    - modify pass 0 of external sort: eliminate unwanted columns
      - read in B pages from E
      - write out  $(T/M) * B$  internally sorted pages of E'
        - refinement: write out  $2*B$  internally sorted pages of E' (on average)
      - tuples in runs - smaller than input tuples
    - modify merging passes: eliminate duplicates
      - number of result tuples is smaller than number of input tuples

## Projection Based on Sorting

- improvement

- \* example

- pass 0:

- scan Exams: 1000 I/Os

- write out 250 pages:

- 20 available buffer pages

- 250 pages => 7 sorted runs about 40 pages long (except the last one, which is about 10 pages long)

- pass 1:

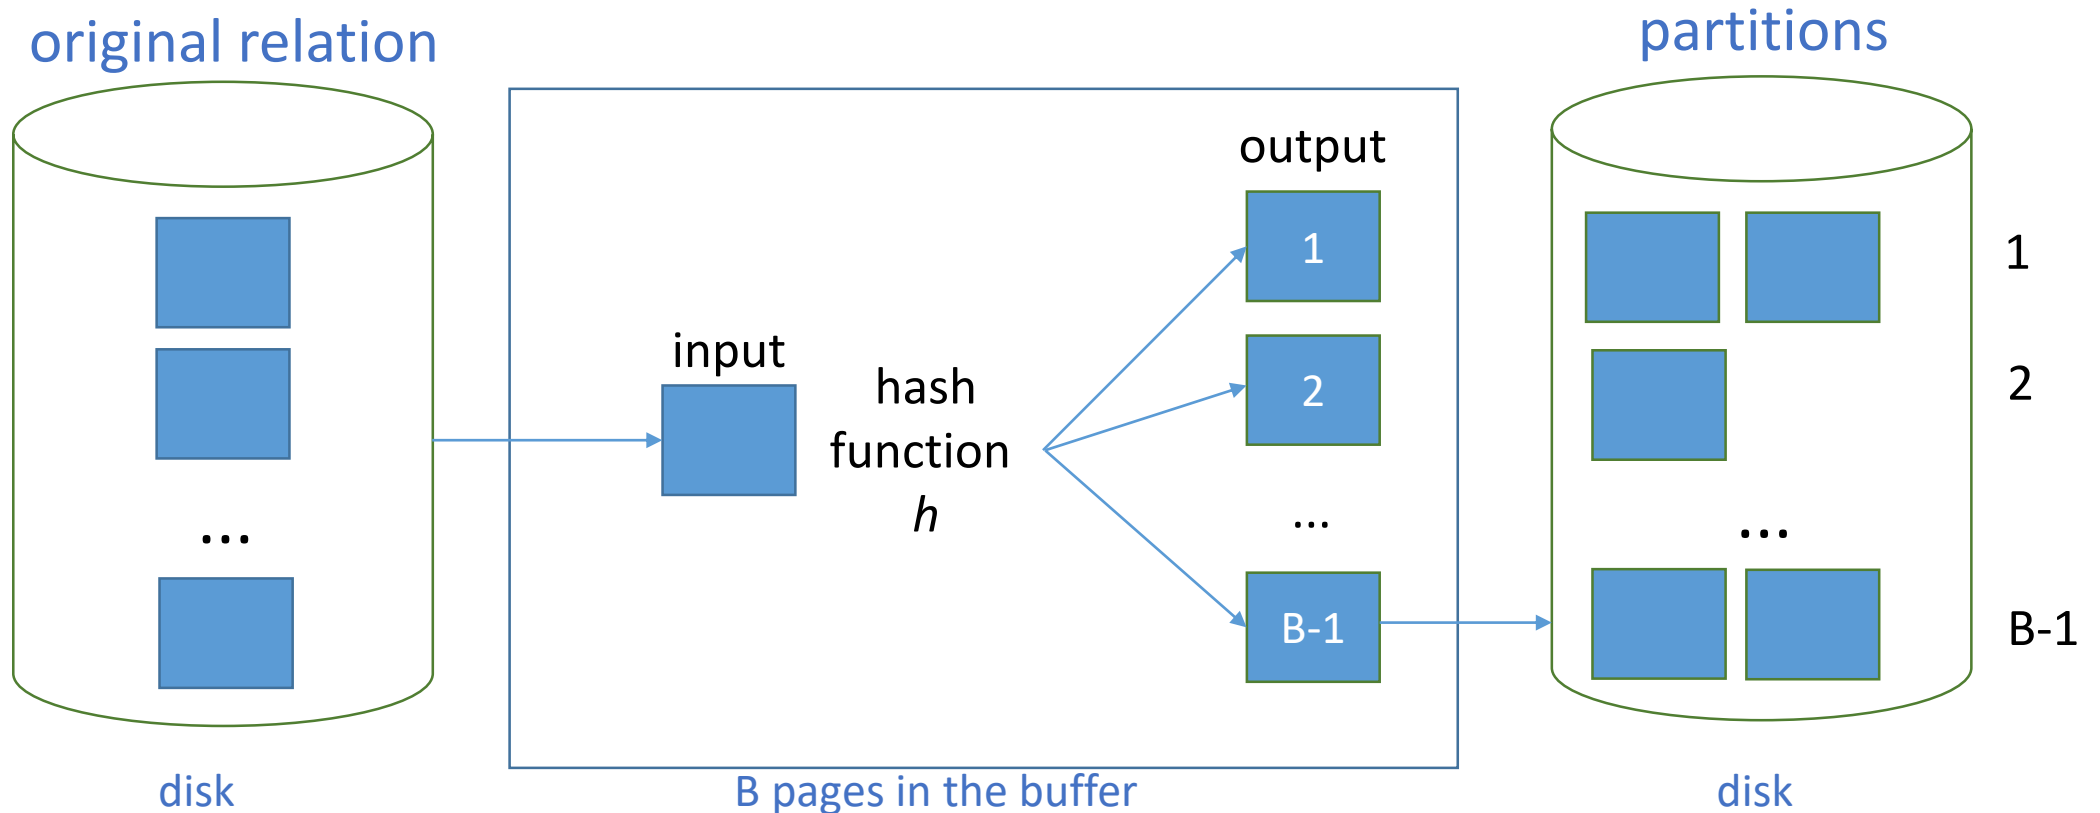
- read in all runs – cost: 250 I/Os

- merge runs

- total cost :  $1000 + 250 + 250 = 1500$  I/Os

## Projection Based on Hashing

- phases: partitioning & duplicate elimination
- partitioning phase:
  - 1 input buffer page – read in the relation one page at a time
  - hash function  $h$  – distribute tuples uniformly to one of  $B-1$  partitions
  - $B-1$  output buffer pages – one output page / partition





## Projection Based on Hashing

- partitioning phase:
  - read the relation using the input buffer page
  - for each tuple  $t$ :
    - discard unwanted fields  $\Rightarrow$  tuple  $t'$
    - apply hash function  $h$  to  $t'$
    - write  $t'$  to the output buffer page that it is hashed to by  $h$

$\Rightarrow$  B-1 partitions

- partition:
  - collection of tuples with:
    - common hash value
    - no unwanted fields
- 2 tuples in different partitions are guaranteed to be distinct

## Projection Based on Hashing

- duplicate elimination phase:
  - process all partitions:
    - read in partition P, one page at a time
      - build in-memory hash table with hash function  $h_2 (\neq h)$  on all fields:
        - if a new tuple hashes to the same value as an existing tuple, compare them to check if they are distinct
        - eliminate duplicates
      - write duplicate-free hash table to result file
      - clear in-memory hash table
- partition overflow
  - apply hash-based projection technique recursively (subpartitions)

## Projection Based on Hashing

- cost
  - partitioning:
    - read E: M I/Os
    - write E': T I/Os
  - duplicate elimination:
    - read in partitions: T I/Os

=> total cost:  $M + 2 * T$  I/Os
- Exams:
  - $1000 + 2 * 250 = 1500$  I/Os

## Set Operations

- intersection, cross-product
  - special cases of join (i.e., join condition for intersection - equality on all fields, no join condition for cross-product)
- union, set-difference
  - similar
- union:  $R \cup S$ 
  - sorting
    - sort R and S on all attributes
    - scan the sorted relations in parallel; merge them, eliminating duplicates
    - refinement
      - produce sorted runs of R and S, merge runs in parallel

## Set Operations

- union:  $R \cup S$ 
  - hashing
    - partition  $R$  and  $S$  with the same hash function  $h$
    - for each  $S$ -partition
      - build in-memory hash table (using  $h_2$ ) for the  $S$ -partition
      - scan corresponding  $R$ -partition, add tuples to hash table, discard duplicates
      - write out hash table
      - clear hash table

## Aggregate Operations

- without grouping
  - scan relation
  - maintain *running information* about scanned tuples
    - COUNT - count of values retrieved
    - SUM - *total* of values retrieved
    - AVG -  $\langle total, count \rangle$  of values retrieved
    - MIN, MAX - smallest / largest value retrieved
- with grouping
  - sort relation on the grouping attributes
  - scan relation to compute aggregate operations for each group
  - improvement: combine sorting with aggregation computation
  - alternative approach based on hashing

## Aggregate Operations

- using existing indexes
  - index with a search key that includes all the attributes required by the query
    - index-only scan
  - attribute list in the GROUP BY clause is a prefix of the index search key (tree index)
    - get data entries (and records, if necessary) in the required order
    - i.e., avoid sorting

# References

- [Ra00] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (2<sup>nd</sup> Edition), McGraw-Hill, 2000
- [Da03] DATE, C.J., An Introduction to Database Systems (8<sup>th</sup> Edition), Addison-Wesley, 2003
- [Ga08] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book, Prentice Hall Press, 2008
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007, <http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si10] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, McGraw-Hill, 2010, <http://codex.cs.yale.edu/avi/db-book/>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems, <http://infolab.stanford.edu/~ullman/fcdb.html>